Page 1: Deep Learning and Transformer Architectures

The Transformer is a neural network architecture based on multi-head self-attention. Unlike recurrent neural networks, Transformers process all tokens simultaneously. Self-attention computes dynamic weights between words regardless of positional distance.

Meeting Notes - Project Apollo

Budget approved: 45000 USD for hardware testing.
Deadline for Phase 2 prototype is November 15th.
Key contact: Dr. Sarah Jenkins in Seattle office.
Action item: Order 12 sensor modules before Friday.